Image generation model comparison

Creating Art

THEN

NOW

Problem Statement

Which deep learning model is the best for image generation?

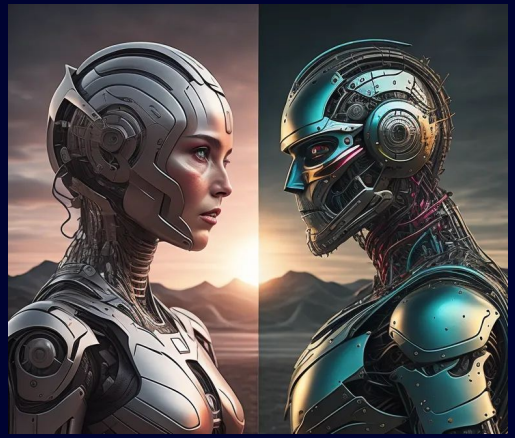

CUB 200-2011

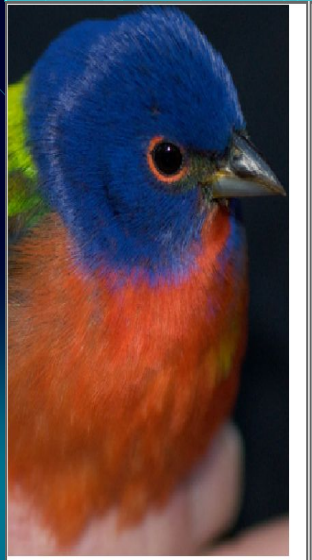

Can't see the bird in the image?
Click here to skip it.
ONLY skip if no bird is visible at all in the image.

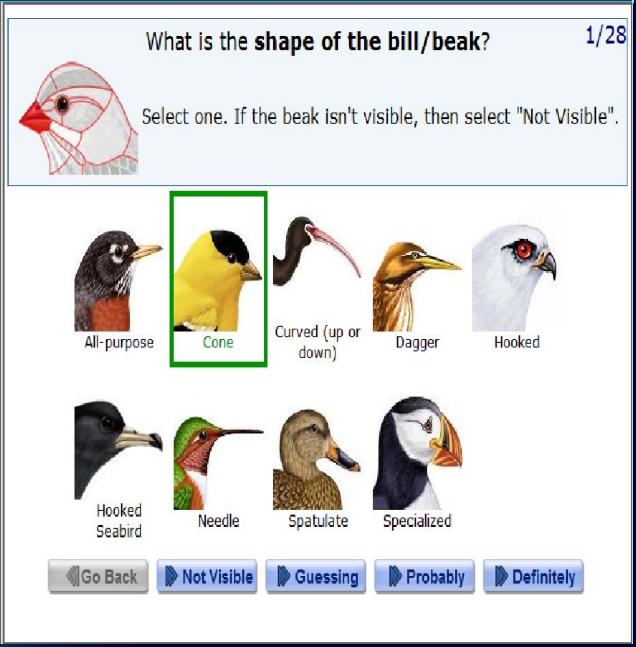

COCO

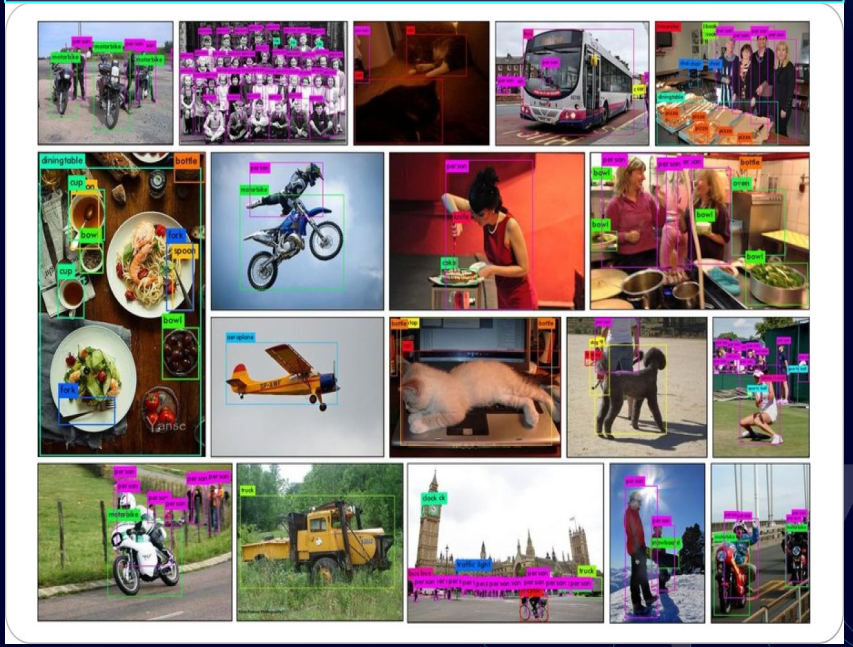

taoxugit/**AttnGAN**

2 Contributors | 78 Issues | 1k Stars | 414 Forks

01

AttnGAN-Attentional Generative Adversarial Networks

How the model work?

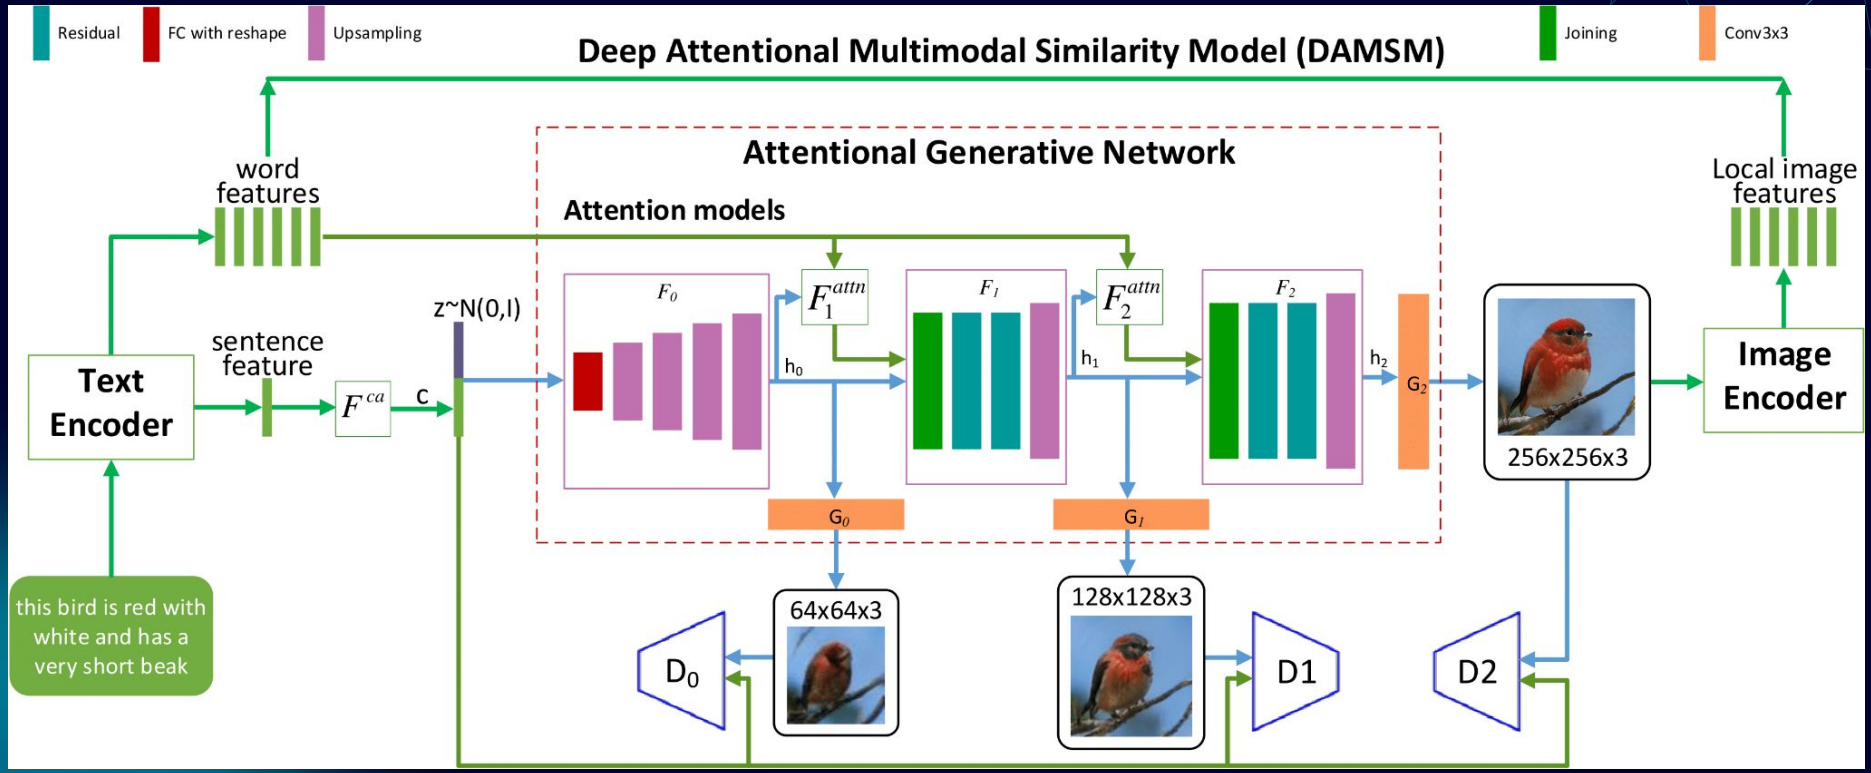

Match keywords with sub-region

Upscaling

ADVANTAGE

01 Produce detailed result

02 Fine-Grained Control

03 Improve text-image matching

tobran/DF-GAN

A Simple and Effective Baseline for Text-to-Image Synthesis (CVPR2022 oral)

2 Contributors | 13 Issues | 281 Stars | 67 Forks

02

DF-GAN: Deep Fusion Generative Adversarial Networks

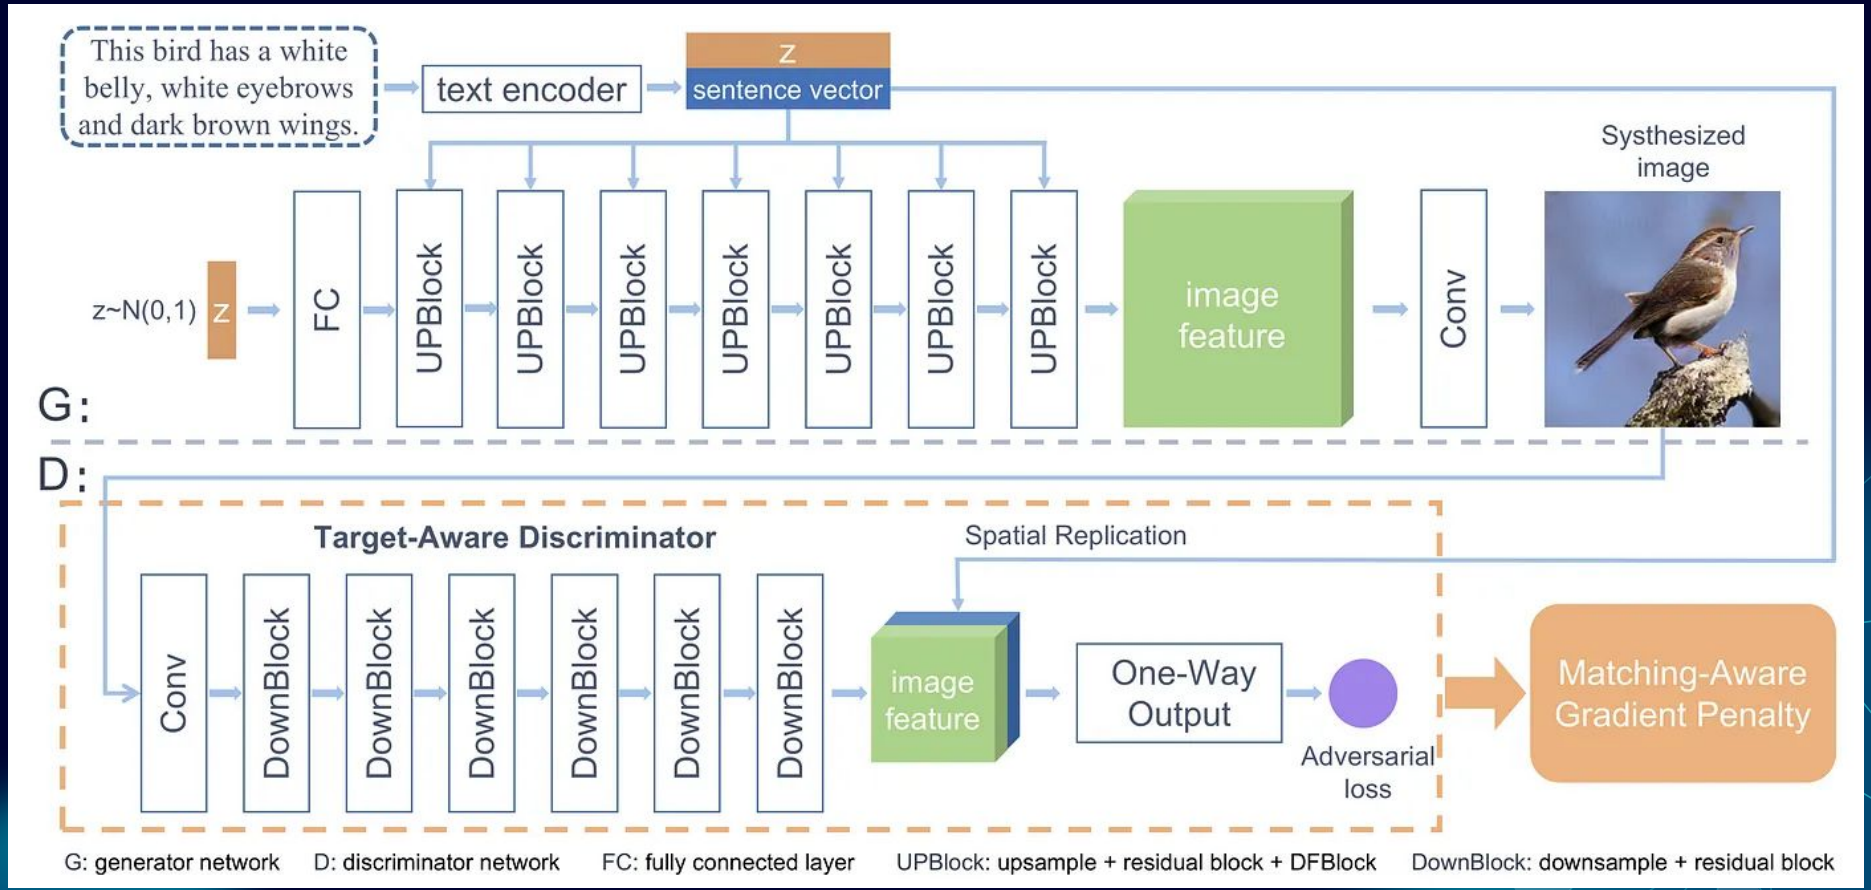
This bird has a white belly, white eyebrows and dark brown wings.
text encoder
z
sentence vector
Systhesized image
z~N(0,1)
z
FC
UPBlock
UPBlock
UPBlock
UPBlock
UPBlock
UPBlock
UPBlock
image feature
Conv
G:
D:
Target-Aware Discriminator
Spatial Replication
Conv
DownBlock
DownBlock
DownBlock
DownBlock
DownBlock
DownBlock
image feature
One-Way Output
Adversarial loss
Matching-Aware Gradient Penalty
G: generator network
D: discriminator network
FC: fully connected layer
UPBlock: upsample + residual block + DFBlock
DownBlock: downsample + residual block

One-Stage Text-to-Image Backbone

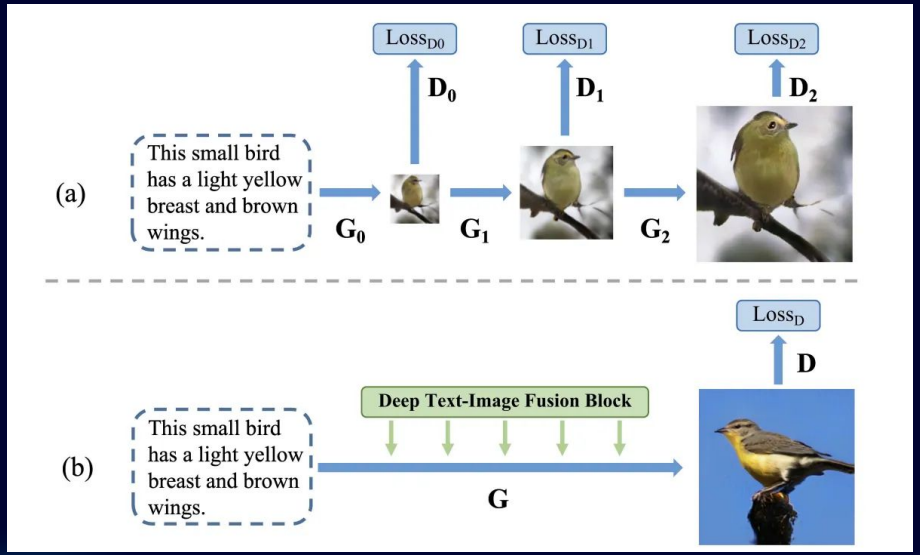

- DF-GAN directly synthesizes high-resolution images from textual descriptions in a single step.
- DF-GAN employs hinge loss to stabilize the adversarial training process.

The Target-Aware Discriminator

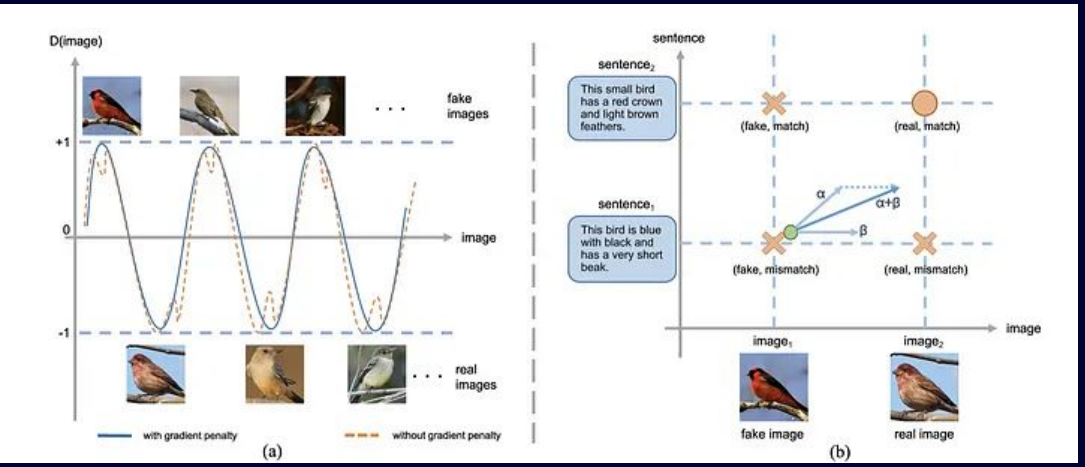

- Introduce Gradient penalty to smoothens the surface for and around the real data points for smoother convergence
- Push the real and text-image consistent points to the minimum of the loss curve

Deep Text-Image Fusion Block (DFBlock)

- In typical text-to-image setups, this blending might be insufficient, resulting in images lacking detail or not matching the text accurately.
- The DFBlock addresses this by employing multiple layers to blend text and visual features more comprehensively.

GALIP

03

GALIP: Generative Adversarial CLIPs

CLIP (Contrastive Language-Image Pre-training)

Neural Network Model developed by OpenAI

- Generate Image from Textual Description
- Image classification

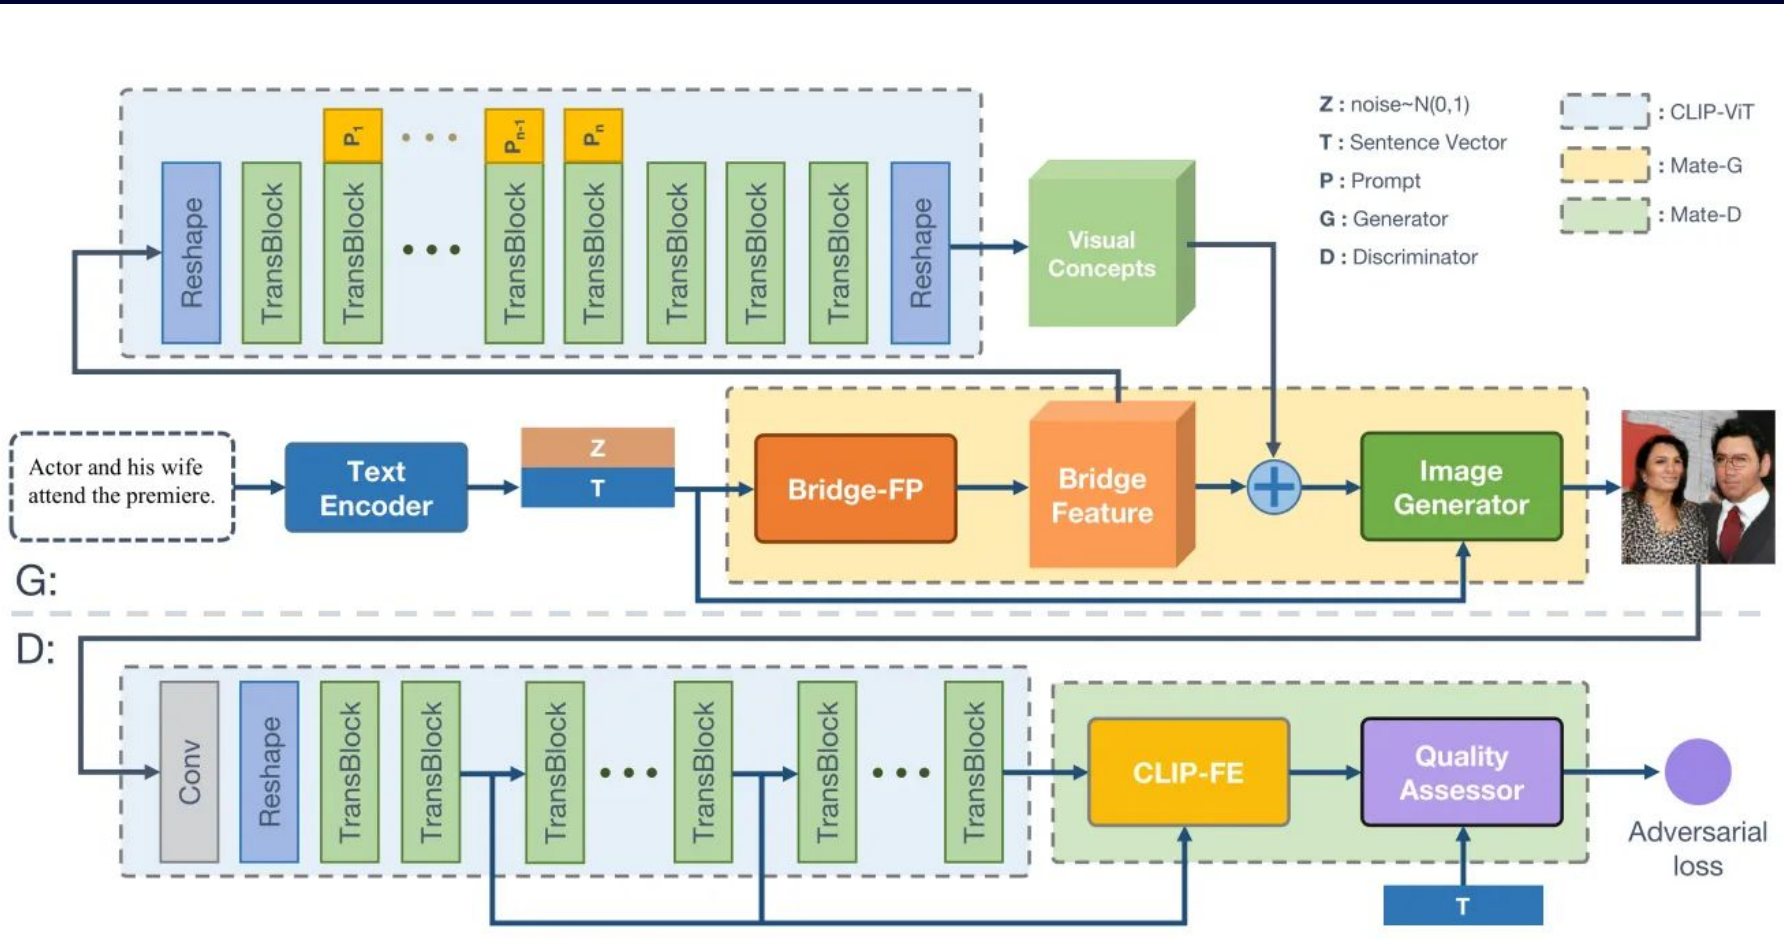
Z : noise~N(0,1)
T : Sentence Vector
P : Prompt
G : Generator
D : Discriminator
: CLIP-ViT
: Mate-G
: Mate-D
Reshape
TransBlock
P1
Pn-1
Pn
Visual Concepts
Actor and his wife attend the premiere.
Text Encoder
Z
T
Bridge-FP
Bridge Feature
Image Generator
G:
D:
Conv
CLIP-FE
Quality Assessor
Adversarial loss

Sample result from each model

CAPTION “The skiers are standing next to a large crowd.”

AttnGAN

DF-GAN

GALIP

Sample result from each model

CAPTION “this bird had brown primaries, a brown crown, and white belly.”

AttnGAN

DF-GAN

GALIP

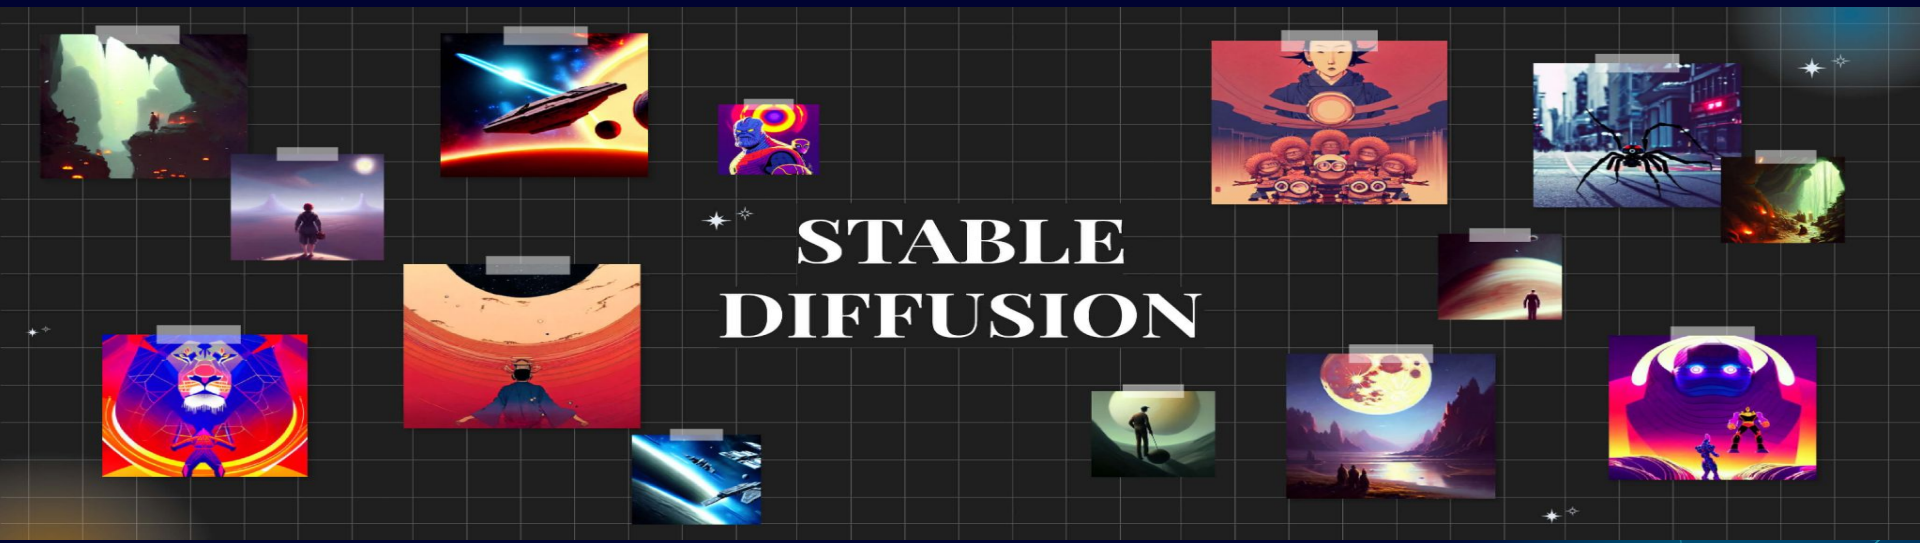

04 Stable Diffusion

Components

Auto-Encoder

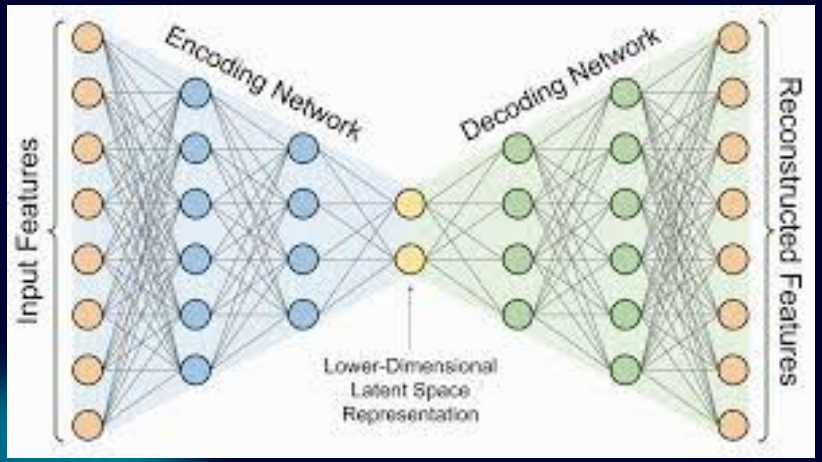

Text-Encoder

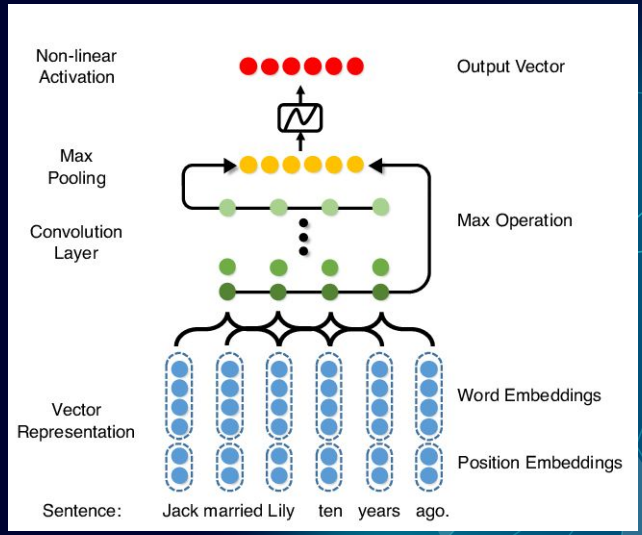

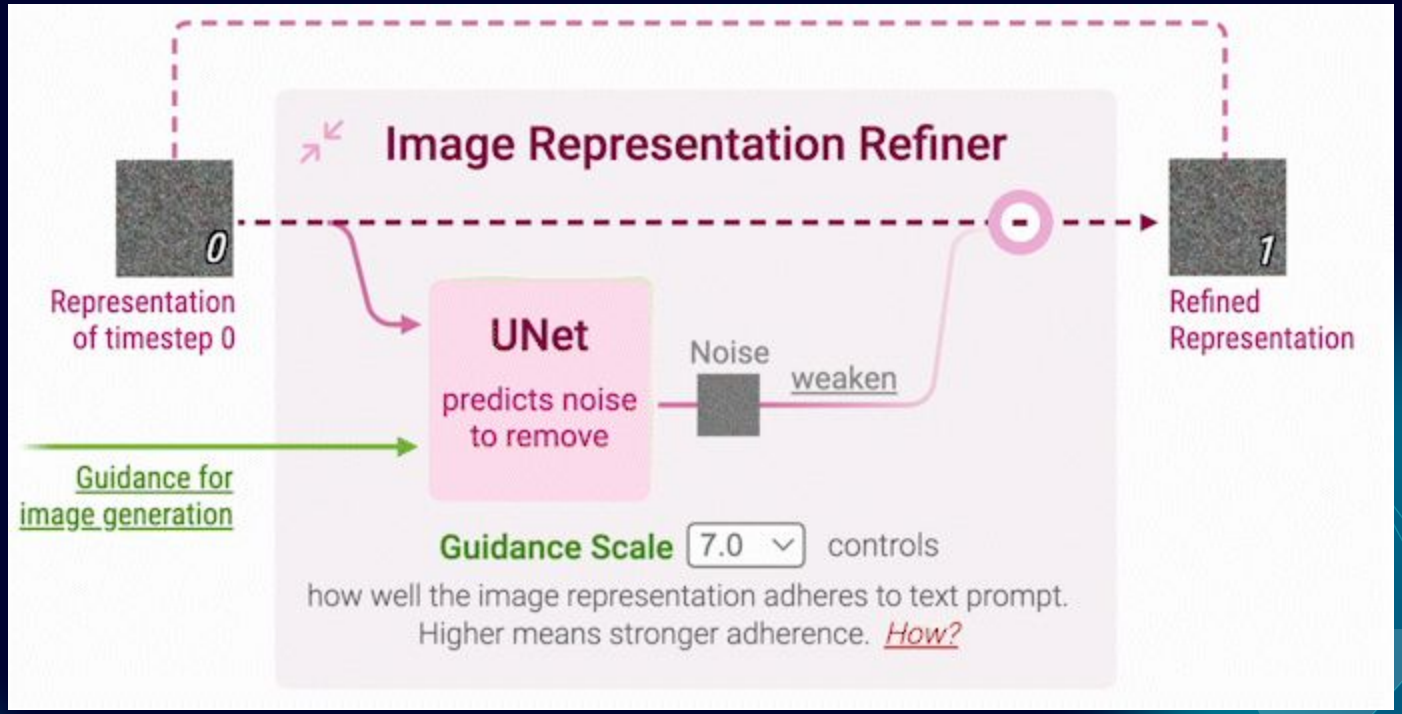
Image Representation Refiner
0
Representation of timestep 0
UNet
predicts noise to remove
Noise
weaken
1
Refined Representation
Guidance for image generation
Guidance Scale
7.0
controls
how well the image representation adheres to text prompt.
Higher means stronger adherence.
How?

santa clause eating a chocolate chip cookie in front of a cozy fireplace
Negative prompt: ugly, bad anatomy, bad proportions, deformed, extra limbs, low quality, low res, mutated, missing limbs, disfigured, disgusting
Steps: 70, Sampler: DPM++ 2M Karras, CFG scale: 7, Seed: 422078003, Size: 512x512, Model hash: 6ce0161689, Model: v1-5-pruned-emaonly, Version: v1.8.0

santa clause eating a chocolate chip cookie in front of a cozy fireplace
Negative prompt: ugly, bad anatomy, bad proportions, deformed, extra limbs, low quality, low res, mutated, missing limbs, disfigured, disgusting
Steps: 70, Sampler: DPM++ 2M Karras, CFG scale: 7, Seed: 3889559361, Size: 512x512, Model hash: 93ed864a22, Model: cyberrealistic_v42, VAE hash: c6a580b13a, VAE: vae-ft-mse-840000-ema-pruned.ckpt, Version: v1.8.0

this bird had brown primaries, a brown crown, and white belly
Negative prompt: ugly, deformed, extra limbs, low quality, low res, mutated, missing limbs, disfigured, disgusting
Steps: 70, Sampler: DPM++ 2M Karras, CFG scale: 7, Seed: 1085542589, Size: 512x512, Model hash: 6ce0161689, Model: v1-5-pruned-emaonly, Version: v1.8.0

this bird had brown primaries, a brown crown, and white belly
Negative prompt: ugly, deformed, extra limbs, low quality, low res, mutated, missing limbs, disfigured, disgusting
Steps: 70, Sampler: DPM++ 2M Karras, CFG scale: 7, Seed: 250506351, Size: 512x512, Model hash: 93ed864a22, Model: cyberrealistic_v42, VAE hash: c6a580b13a, VAE: vae-ft-mse-840000-ema-pruned.ckpt, Version: v1.8.0

The skiers are standing next to a large crowd
Negative prompt: ugly, deformed, extra limbs, low quality, low res, mutated, missing limbs, disfigured, disgusting

Steps: 70, Sampler: DPM++ 2M Karras, CFG scale: 7, Seed: 1179824724, Size: 512x512, Model hash: 6ce0161689, Model: v1-5-pruned-emaonly, Version: v1.8.0

The skiers are standing next to a large crowd
Negative prompt: ugly, deformed, extra limbs, low quality, low res, mutated, missing limbs, disfigured, disgusting

Steps: 70, Sampler: DPM++ 2M Karras, CFG scale: 7, Seed: 435320631, Size: 512x512, Model hash: 93ed864a22, Model: cyberrealistic_v42, VAE hash: c6a580b13a, VAE: vae-ft-mse-840000-ema-pruned.ckpt, Version: v1.8.0

Diffusion model pros and cons

- Enhancement from low-quality source
- Feature- specific enhancement
- open-sourced

- Require lot of training data
- Time Consuming
- Synthesized visual features
- Computationally intensive

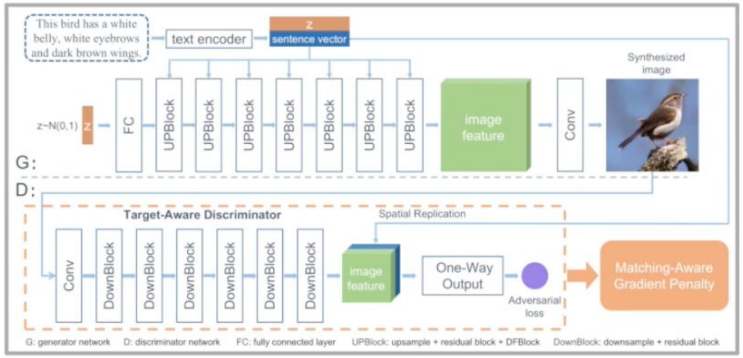

The framework of DF-GAN [1]

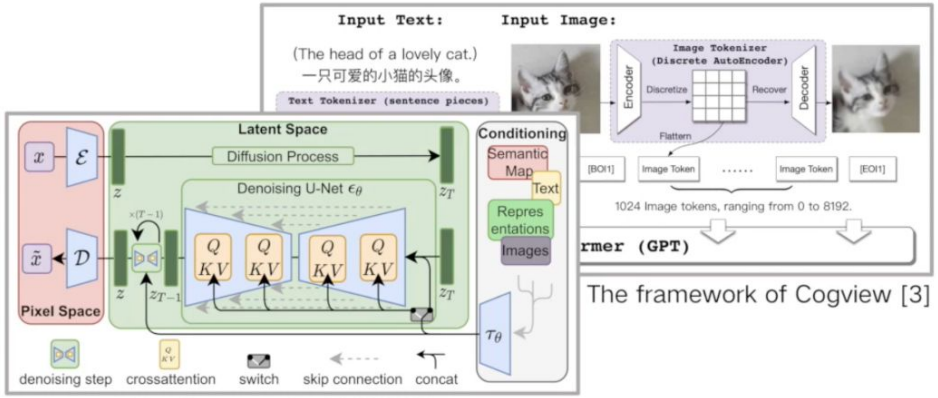

The framework of Cogview [3]

The framework of LDM [2]

☐ GAN

- hard to synthesize complex images
- ✓ fast synthesis speed
- ✓ small model size
- ✓ meaningful latent space

☐ AR and diffusion models

- ✓ more powerful generative capabilities
- slow synthesis speed
- large model size and hardware requirements
- lack a meaningful latent space

THANK YOU!

Resources

https://www.vision.caltech.edu/datasets/cub_200_2011/
https://paperswithcode.com/dataset/cub-200-2011
https://codeburst.io/understanding-attngan-text-to-image-convertor-a79f415a4e89
https://blog.segmind.com/stable-diffusion-deployment/#the-anatomy-of-stable-diffusion-
https://poloclub.github.io/diffusion-explainer/#:~:text=Stable%20Diffusion%20generates%20an%20image,quality%20of%20the%20image%20representation.

Citations

AttnGAN:

```
@article{Tao18attngan,
        author    = {Tao Xu, Pengchuan Zhang, Qiuyuan Huang, Han Zhang, Zhe Gan, Xiaolei
                     Huang, Xiaodong He},
        title     = {AttnGAN: Fine-Grained Text to Image Generation with
                     Attentional Generative Adversarial Networks},
        Year = {2018},
        booktitle = {{CVPR}}
       }
```

Citations

DF-GAN:

```
@inproceedings{tao2022df,
        title={DF-GAN: A Simple and Effective Baseline for Text-to-Image Synthesis},
        author={Tao, Ming and Tang, Hao and Wu, Fei and Jing, Xiao-Yuan and Bao,
        Bing-Kun and Xu, Changsheng},
        booktitle={Proceedings of the IEEE/CVF Conference on Computer Vision
                and Pattern Recognition},pages={16515--16525}, year={2022}
        }
```

Citations

GALIP:

```
@inproceedings{tao2023galip,
        title={GALIP: Generative Adversarial CLIPs for Text-to-Image Synthesis},
        author={Tao, Ming and Bao, Bing-Kun and Tang, Hao and Xu, Changsheng},
        booktitle={Proceedings of the IEEE/CVF Conference on Computer Vision
                and Pattern Recognition},
        pages={14214--14223},
        year={2023}
       }
```

Citations

Stable Diffusion

```
authors:
    - given-names: AUTOMATIC1111
    title: "Stable Diffusion Web UI"
    date-released: 2022-08-22
    url: "https://github.com/AUTOMATIC1111/stable-diffusion-webui"
```